

RaaSI: Developer Guide

Michael Coffey

July 5, 2022

Contents

1	Introduction	1
2	Installation	1
2.1	RaaSI Setup	2
2.2	Secondary Service Stateful Storage	3
3	Configuration	4
3.1	Primary Service Monitoring API	5
3.2	Kubernetes Configuration	6
3.3	Developing Secondary Services	7
4	Conclusion	7

1 Introduction

The purpose of this document is to give developers who are using RaaSI a more technical look at how the system can be used and customized. For a more generic view of the use cases of RaaSI (including how to get RaaSI installed and configured), users can view the "RaaSI Users Guide." This developers guide will show users how to customize their installation of RaaSI through advanced manipulation, configure their RaaSI cluster to best match their needs, instruct the user on how to develop and manage secondary services used in RaaSI, and give an insight on how the source code of RaaSI is managed.

2 Installation

The basic installation of RaaSI can be completed through the RaaSI CLI and by following the steps detailed in the Users Guide. This installation process can be customized to more closely fit the needs of users by the modification of scripts found in the RaaSI CLI. These scripts are written in bash to allow users to customize them in any manor that best fits their specific situation. The scripts should be seen as a guide to RaaSI rather than RaaSI itself. Through script

modification, users can easily create more load balancers and stateful storage nodes, as well as change the type of stateful storage that secondary services use.

2.1 RaaSI Setup

The RaaSI CLI guides users through the installation of all components used in RaaSI; however, it limits the users to have only three load balancers and three storage nodes. The load balancer and storage node count is set to three by default because three nodes are the minimum amount needed to provide accurate voting between the nodes to make communication decisions. This number can be increased if the user has available nodes that they would like to use for the purpose of load balancing or storing the stateful data of secondary services.

In order to change the number of load balancers used in RaaSI, developers should navigate to the "LoadBalancerSetup/LoadBalancerInstallation.sh" script shown in figure 1. Once in the file, the developer can add as many load balancers to the installation script as they desire. Due to all of the load balancers using a virtual IP (VIP) to communicate between Master and Worker Nodes, this is the only script that needs modification.

Figure 1: RaaSI Load Balancer Installation Script

```
1 #!/bin/bash
2
3 echo "Beginning Load Balancer setup script..."
4
5 echo "Enter the ip address for LoadBalancer1: "
6 read -r LB1_ip
7
8 # validation
9 val=$(./Validation/checkValidation.sh "$LB1_ip" 1)
10 while [ "$passed" != "$val" ];
11 do
12     echo "Unexpected Response: expected IP address"
13     echo "Enter the ip address for LoadBalancer1: "
14     read -r LB1_ip
15     # validation
16     val=$(./Validation/checkValidation.sh "$LB1_ip" 1)
17 done
18
19 echo "Enter the ip address for LoadBalancer2: "
20 read -r LB2_ip
21
22 # validation
23 val=$(./Validation/checkValidation.sh "$LB2_ip" 1)
24 while [ "$passed" != "$val" ];
25 do
26     echo "Unexpected Response: expected IP address"
27     echo "Enter the ip address for LoadBalancer2: "
28     read -r LB2_ip
29     # validation
30     val=$(./Validation/checkValidation.sh "$LB2_ip" 1)
31 done
32
33 echo "Enter the ip address for LoadBalancer3: "
34 read -r LB3_ip
35
36 # validation
37 val=$(./Validation/checkValidation.sh "$LB3_ip" 1)
38 while [ "$passed" != "$val" ];
39 do
40     echo "Unexpected Response: expected IP address"
41     echo "Enter the ip address for LoadBalancer3: "
42     read -r LB3_ip
43     # validation
44     val=$(./Validation/checkValidation.sh "$LB3_ip" 1)
45 done
46
```

It is a slightly more complex process if developers want to change to number of secondary service storage nodes due to the amount of scripts that need to be changed for this to work. The first script that developers will need to change to modify the count of storage nodes is "GlusterSetup/ GlusterInstal-

lationMaster.sh” shown in figure 2. To add more storage nodes to this script, developers need to request more IP address and login credentials from the user of the script. Developers also must modify the Worker Node installation

Figure 2: RaaSI Storage Node Installation Script

```

1 #!/bin/bash
2
3 echo "Beginning GlusterFS setup script..."
4
5 echo "Enter the ip address for Storage1: "
6 read -r storage1_ip
7
8 # validation
9 val=$(../Validation/checkValidation.sh "$storage1_ip" 1)
10 while [ "passed" != "$val" ];
11 do
12     echo "Unexpected Response: expected IP address"
13     echo "Enter the ip address for Storage1: "
14     read -r storage1_ip
15     # validation
16     val=$(../Validation/checkValidation.sh "$storage1_ip" 1)
17 done
18
19 echo "Enter the ip address for Storage2: "
20 read -r storage2_ip
21
22 # validation
23 val=$(../Validation/checkValidation.sh "$storage2_ip" 1)
24 while [ "passed" != "$val" ];
25 do
26     echo "Unexpected Response: expected IP address"
27     echo "Enter the ip address for Storage2: "
28     read -r storage2_ip
29     # validation
30     val=$(../Validation/checkValidation.sh "$storage2_ip" 1)
31 done
32
33 echo "Enter the ip address for Storage3: "
34 read -r storage3_ip
35
36 # validation
37 val=$(../Validation/checkValidation.sh "$storage3_ip" 1)
38 while [ "passed" != "$val" ];
39 do
40     echo "Unexpected Response: expected IP address"
41     echo "Enter the ip address for Storage3: "
42     read -r storage3_ip
43     # validation
44     val=$(../Validation/checkValidation.sh "$storage3_ip" 1)
45 done
46

```

script, found at "ClientNodeSetup/ glusterfsClientSetup.sh," to make sure that the Worker Nodes are connected to the newly added storage nodes. This script is shown in figure 3. Similar to the "GlusterInstallationMaster.sh" script, this script just needs to have the extra IP address and login information added to it to allow developers to increase the number of storage nodes.

2.2 Secondary Service Stateful Storage

For secondary service stateful storage RaaSI uses GlusterFS. The reason behind this is that GlusterFS is a lightweight, easy to integrate, file system that allows users to manage with speed and simplicity. However, this does not have to be the storage used for secondary services in RaaSI. Developers can disregard GlusterFs and use their own form of storage if they so choose.

Modification of the type of stateful storage for secondary services can be done by modifying the script "GlusterSetup/ GlusterInstallationMaster.sh," shown in figure 2, and replacing the installation of GlusterFS information with the installation process of the storage that the developer desires. If developers choose to change this storage, they will also need to modify the scripts "ClientNodeSetup/ glusterfsClientSetup.sh", shown in figure 3, and "ClientNodeSetup/ clientSideNodeSetup.sh", shown in figure 4. The script "clientSideNodeSetup.sh" is where

Figure 3: RaaSI Worker Node Storage Installation Script

```
1 #!/bin/bash
2
3 client_ip=$1
4 client_user=$2
5
6 echo "Connecting the client node to GlusterFS..."
7
8 echo "Enter the IP address for Storage1:"
9 read -r storage1_ip
10
11 # validation
12 val=$(../Validation/checkValidation.sh "$storage1_ip" 1)
13 while [ "passed" != "$val" ];
14 do
15     echo "Unexpected Response: expected IP address"
16     echo "Enter the IP address for Storage1: "
17     read -r storage1_ip
18     # validation
19     val=$(../Validation/checkValidation.sh "$storage1_ip" 1)
20 done
21
22 echo "Enter username to connect to on Storage1: "
23 read -r storage_user1
24
25 # validation
26 val=$(../Validation/checkValidation.sh "$storage_user1" 2)
27 while [ "passed" != "$val" ];
28 do
29     echo "Unexpected Response: valid username was expected"
30     echo "Enter username to connect to on Storage1: "
31     read -r storage_user1
32     # validation
33     val=$(../Validation/checkValidation.sh "$storage_user1" 2)
34 done
35
36 echo "Enter the IP address for Storage2:"
37 read -r storage2_ip
38
39 # validation
40 val=$(../Validation/checkValidation.sh "$storage2_ip" 1)
41 while [ "passed" != "$val" ];
42 do
43     echo "Unexpected Response: expected IP address"
44     echo "Enter the IP address for Storage2: "
45     read -r storage2_ip
46     # validation
47     val=$(../Validation/checkValidation.sh "$storage2_ip" 1)
48 done
49
50 echo "Enter username to connect to on Storage2: "
51 read -r storage_user2
```

GlusterFS is first installed to the Worker Nodes, this script must be modified to install the new storage to the Worker Nodes instead. The script "glusterfsClientSetup.sh" is where GlusterFS is configured on the Worker Nodes. This is where developers should place the configuration steps for the storage that they choose to use.

3 Configuration

Once RaaSI has been installed, developers can configure the environment to better fit their use. As discussed in the Users Guide, the Primary Service Monitoring API of RaaSI can be configured to monitor the health of any primary service through modifying the python script for monitoring services. RaaSI can also be configured through changing the dashboard and overlay network that is setup in the kubernetes configuration. Finally, secondary services can be managed from their creation and throughout their deployment with the use of the stateful storage system, kubernetes node communication, or by creating secondary services that support other forms of communication.

Figure 4: RaaS Worker Node Installation Script

```

10 echo "Have you already completed Kubernetes Installation(Y/N)"
11 read -r response
12
13 # validation
14 val=$(./Validation/checkValidation.sh "$response" 0)
15 while [ "$val" != "passed" ]; do
16     echo "Unexpected Response"
17     echo "Have you already completed Kubernetes Installation(Y/N)"
18     read -r response
19     # validation
20     val=$(./Validation/checkValidation.sh "$response" 0)
21 done
22
23 if [[ "$response" == "y" || "$response" == "Y" ]]; then
24     echo "Kubernetes Installation Complete."
25 else
26     echo "Beginning Kubernetes Installation"
27     apt-get install -y openssh-server
28     apt-get install -y docker.io
29     apt-get install -y apt-transport-https curl
30     apt-get install -y wget
31     wget https://download.gluster.org/pub/gluster/glusterfs/7/raa.pub | apt-key add -
32     echo "deb [arch=amd64] https://download.gluster.org/pub/gluster/glusterfs/7/LATEST/debian/buster/amd64/apt-buster main" > /etc/apt/sources.list.d/gluster.list
33
34     curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
35     cat <EOF> > /etc/apt/sources.list.d/kubernetes.list
36     deb http://apt.kubernetes.io/ kubernetes-xenial main
37 EOF
38
39 apt-get update
40 apt-get install -y glusterfs-client
41
42 apt-get install -y kubelet kubeadm kubectl
43
44 echo "Kubernetes Installation Complete."
45 fi
46
47 echo "Have you already completed Kubernetes Configuration(Y/N)"
48 read -r response
49

```

3.1 Primary Service Monitoring API

To monitor a new primary service with the primary service monitoring API, developers must add the primary service to the database used by the API, add the method used to monitor the health of the service, and add the new service to "all-services.yaml" file. First, the developer should check the "PrimaryServiceMonitoringAPI/ data/ schema.sql" file to see if their service name is already in the database. If it is not, the developer should add the new service to the bottom of the file in the form of "INSERT INTO service_type(service_type_name) VALUES('< NewService >')" (replacing < NewService > with the name of the service to add. This file is shown in figure 5.

After the new service has been added to the database, the developer should add a health monitoring method for the service. This can be done in the "PrimaryServiceMonitoringAPI/ judge/ tasks.py" file. In this file the developer should create a new method that checks the health of the service that they are adding. Examples of these methods are shown throughout this file such as monitoring a web server shown in figure 6. Once this method has been created, the developer should add the reference to their new service and the health monitoring method that is used for the service to the "poll()" method as shown in figure 7.

Finally, the developer should navigate to the "PrimaryServiceMonitoringAPI/ all-services.yaml" file and add their new primary service to the file under the "services" section of the file. In this file, the variable "service_type_name" is the name of the service type that you added to the database, the variable "service_name" is a reference name to the primary service (it's important to give

Figure 5: Primary Service Monitor API Schema

```

Open  [R] schema.sql  Save  -  ⌵  ⌵
~/RaaS/RaaS-main/PrimaryServiceMonitor/API/data

34 FOREIGN KEY (service_id) REFERENCES service(service_id),
35 -- FOREIGN KEY (team_id) REFERENCES team(team_id),
36 FOREIGN KEY (service_type_name) REFERENCES service_type(service_type_name)
37 );
38
39 CREATE TABLE error (
40   error_id INTEGER PRIMARY KEY ASC,
41   service_id INTEGER NOT NULL,
42   error_message TEXT NOT NULL,
43   error_time DATETIME DEFAULT CURRENT_TIMESTAMP,
44   FOREIGN KEY (service_id) REFERENCES service(service_id)
45 );
46
47 CREATE TABLE deployment (
48   deployment_id INTEGER PRIMARY KEY ASC,
49   deployment_name TEXT NOT NULL,
50   deployment_device TEXT NOT NULL,
51   deployment_percentage INTEGER NOT NULL,
52   service_name TEXT NOT NULL,
53   service_id INTEGER NOT NULL,
54   FOREIGN KEY (service_id) REFERENCES service(service_id)
55 );
56
57 INSERT INTO service_type(service_type_name) VALUES('dns');
58 INSERT INTO service_type(service_type_name) VALUES('http');
59 INSERT INTO service_type(service_type_name) VALUES('https');
60 INSERT INTO service_type(service_type_name) VALUES('ftp');
61 INSERT INTO service_type(service_type_name) VALUES('nfs');
62 INSERT INTO service_type(service_type_name) VALUES('ssh');
63 INSERT INTO service_type(service_type_name) VALUES('ftp');
64 INSERT INTO service_type(service_type_name) VALUES('mail');
65 INSERT INTO service_type(service_type_name) VALUES('db');
66 INSERT INTO service_type(service_type_name) VALUES('mysql');
67 INSERT INTO service_type(service_type_name) VALUES('rdp');
68 INSERT INTO service_type(service_type_name) VALUES('pop');
69 INSERT INTO service_type(service_type_name) VALUES('snb');
70 INSERT INTO service_type(service_type_name) VALUES('ad');

```


this a unique name if there are multiple primary services of the same type), the variable "service_connection" is the location of this primary service (this can be an IP address or DNS hostname), and the variable "service_request" is used to add arguments to the health monitor check. The example shown curls google.com and checks if "home" is anywhere in the response.

3.2 Kubernetes Configuration

The dashboard that is currently used for the default RaaSI installation [1] is a default dashboard used to display basic kubernetes metrics. This dashboard, shows the state of the kubernetes cluster; including node health, deployment status, cluster errors, and node resource limitations. However, if developers would prefer to use a different dashboard to show the metrics of the system they need to modify the script "MasterNodeSetup/ masterSetup.sh" shown in figure 8. To use a different dashboard, developers can swap out the current dashboard yaml location with their own custom one.

Developers can also change the type of overlay network that kubernetes uses for node communication. Currently, RaaSI uses the tigra calico overlay network, which integrates nicely with kubernetes. It is simple to change this overlay network, the developer just needs to change the source of the yaml file for the overlay network shown in the "MasterNodeSetup/ masterSetup.sh" script in figure 8. Once this change is complete, the developer can run the "masterSetup.sh" script and the new overlay network should be setup.

Figure 6: Primary Service Monitor API Health Monitor 2



```

108
109
110 @app.task(soft_time_limit=0)
111 def poll_web(poll_timeout, service_id, service_type, service_connection, service_request,
112             service_running, type1):
113     deployments = execute_db_query('select * from deployment where service_id = ?',
114                                   [service_id])
115     try:
116         try:
117             result = s.get(service_type + '://' + service_connection, timeout=poll_timeout,
118                           verify=False).text
119         except requests.exception.Timeout as e:
120             execute_db_query('insert into error(service_id, error_message) values(?,?)',
121                             [service_id, 'HTTP(s) Request resulted in a Timeout exception: ' + repr(e)])
122             if service_running == 1:
123                 # execute deployment
124                 for deployment in deployments:
125                     deploy(deployment)
126             execute_db_query('update service set service_running = 0 where service_id
127                             = ?', [service_id])
128             return
129         except requests.exception.ConnectionError as e:
130             execute_db_query('insert into error(service_id, error_message) values(?,?)',
131                             [service_id, 'HTTP(s) Request resulted in a ConnectionError exception: ' + repr(e)])
132             if service_running == 1:
133                 # execute deployment
134                 for deployment in deployments:
135                     deploy(deployment)
136             execute_db_query('update service set service_running = 0 where service_id
137                             = ?', [service_id])
138             return
139         except requests.exception.HTTPError as e:
140             execute_db_query('insert into error(service_id, error_message) values(?,?)',
141                             [service_id, 'HTTP(s) Request resulted in a HTTPError exception: ' + repr(e)])
142             if service_running == 1:
143                 # execute deployment

```

3.3 Developing Secondary Services

Secondary services can be deployed through the script presented in the RaaSI CLI; however, this deployment is only the beginning of managing the secondary services. Developers can use the stateful storage system to perform communication between secondary service instances. This can be done by passing data through the mounted drive on the secondary service. Developers can also send messages directly to secondary services through the "kubectl exec" command. This command allows developers to send commands to nodes belonging to a specific deployment or to individual nodes (the IP address of all nodes can be shown through viewing the dashboard). These forms of communication allow developers to manipulate their secondary services to best fit their scenarios.

4 Conclusion

This concludes the RaaSI Developers Guide. RaaSI can be found at the following git repository [2]. For more generic information regarding the installation and configuration of RaaSI, users should read the RaaSI Users Guide. There is also a technical report that has been created demonstrating the need for RaaSI, works that are related to RaaSI, and discusses tests run on RaaSI. That report is titled RaaSI Technical Report.

